

EARNINGS TO AI STRATEGIST

From public evidence to one testable AI pilot

A reader-facing guide to the decision workflow, transparent scoring model, 30-day pilot framework, and two public-company demonstrations.

Created by Vedant Jain

A specialized Hermes agent for evidence-backed
AI opportunity discovery.



Watch the demo



Explore the repository

A decision workflow, not another idea generator

The goal is to help an executive decide which operating problem deserves an AI pilot, what evidence supports it, and what would make the team stop.

AI opportunity discovery often breaks down before implementation begins. Evidence is scattered, generic ideation produces long lists, and uncertainty gets hidden behind polished language.

Scattered evidence

Signals sit across earnings releases, SEC filings, presentations, and transcripts.

Weak decision links

Use cases often lack a named user, workflow, decision owner, dataset, or baseline.

Blurred certainty

Management claims, analyst interpretation, and unverified assumptions can be presented as equivalent.

Premature value claims

ROI, data readiness, and feasibility are often implied before internal evidence exists.

Scope: This project addresses the front end of AI strategy work. It creates a defensible opportunity hypothesis and a bounded validation plan. It does not replace internal discovery, security review, process mapping, technical architecture, or financial approval.

Evidence earns the recommendation



1	Company and period Resolve the public company, latest completed fiscal period, and analysis date.
2	Source manifest Prioritize company investor materials and the matching SEC filing.
3	Evidence ledger Verify exact quotations and retain date, location, provenance, and limitations.
4	Three problems Separate facts, inference, assumptions, counter-signals, and missing evidence.
5	Ranked opportunities Tie each option to a user, workflow, decision, required data, and non-AI baseline.
6	One pilot Define baseline, target, review points, guardrails, stop conditions, and day-30 decision.

Transparent scoring that can be challenged

$$\text{Weighted score} = 0.35 \times \text{impact} + 0.25 \times \text{feasibility} + 0.25 \times \text{data readiness} + 0.15 \times (6 - \text{risk})$$

Dimension	Weight	What the reviewer asks
Impact	35%	Would this materially improve a named decision or operating outcome?
Feasibility	25%	Can a bounded test be built and evaluated within the proposed workflow?
Data readiness	25%	Do the required records, labels, permissions, and baselines plausibly exist?
Risk-adjusted safety	15%	Can the test remain reversible, human-controlled, and appropriately governed?

The score is a decision aid, not proof. Public information cannot establish internal data quality, implementation effort, ownership, operating cost, or realized value. Every score carries a rationale and a sensitivity warning. If reasonable changes reverse the winner, the ranking is labeled fragile.

Why visible weights matter

Executives can disagree with the formula and still inspect it. The model does not hide the trade-offs inside an opaque recommendation. A future interface can let reviewers change the weights and immediately see whether the selected pilot changes.

A 30-day pilot built to learn and stop

Days 1-5	Validate	Confirm the workflow, owner, data, baseline, security requirements, and stop authority.
Days 6-12	Build	Create a frozen-data prototype and a reproducible non-AI baseline.
Days 13-21	Test	Evaluate errors, user acceptance, edge cases, and red-team scenarios.
Days 22-30	Shadow	Run beside the existing process with no autonomous high-impact action.

Day 30 ends with a named decision: stop, extend, or proceed to a controlled next phase.

Every pilot package includes:

- One recurring decision and one accountable owner role.
- A frozen non-AI baseline and proposed acceptance thresholds.
- Human approval points, escalation paths, and override logging.
- Data, privacy, security, legal, and model-risk prerequisites.
- Explicit non-goals and immediate stop conditions.
- A measurement plan that avoids extrapolating pilot results into unproven enterprise ROI.

PepsiCo: order-fill exception prioritization

Independent public-data demonstration using Q2 2026 materials. Not affiliated with or endorsed by PepsiCo.

3.75 / 5

top provisional opportunity
score

30 days

shadow-mode evaluation

5 days

data and baseline gate

Evidence set: Q2 2026 earnings release, prepared management remarks, and Form 10-Q.

Operating themes: affordability versus net-revenue quality; service, fill, and unit-cost improvement; productivity execution under cost pressure.

Recommended pilot hypothesis

AI-assisted order-fill exception prioritization in one PepsiCo Foods North America workflow. The tool would rank eligible exceptions and show the source fields behind each recommendation. An authorized planner would retain every operational decision.

- **In scope:** one operating unit, one exception type, one daily queue, and one existing resolution workflow.
- **Human control:** planners accept, reject, or reorder recommendations; existing systems execute any action.
- **Primary measures:** median triage time, high-priority recall before cutoff, and false-priority rate.
- **Boundary:** no autonomous allocation, order modification, pricing, customer communication, or ROI claim.
- **Why the ranking is fragile:** public evidence names priorities but does not prove the existence, volume, root causes, or economic value of order exceptions.

[Open the PepsiCo case-study files on GitHub](#)

JPMorgan Chase: AI portfolio evidence review

Independent public-data demonstration using Q2 2026 materials. Not affiliated with or endorsed by JPMorgan Chase.

4.10 / 5

top provisional opportunity
score

15-25

assumed in-scope use
cases

100%

target traceability for
included claims

Evidence set: Q2 2026 earnings release, investor presentation, company-published transcript, and Form 10-Q.

Operating themes: AI portfolio prioritization, customer-outcome evidence, and human-controlled workforce transition.

Recommended pilot hypothesis

An AI use-case evidence and portfolio review assistant for one non-customer-facing operations portfolio. It would assemble traceable evidence for a monthly human decision to advance, hold, merge, or retire selected use cases.

- **In scope:** read-only use-case records, approved benefit measurements, control status, adoption telemetry, incidents, and decision history.
- **Human control:** authorized reviewers retain every funding, lifecycle, control, deployment, and people decision.
- **Primary measures:** packet-preparation time, evidence completeness, traceability, rework, and actionability of flags.
- **Boundary:** no generated ROI forecast, control waiver, client communication, employment action, or automatic lifecycle update.
- **Built-in stop test:** stop if existing governance already provides timely, complete, evidence-linked reviews at lower complexity.

[Open the JPMorgan Chase case-study files on GitHub](#)

The strongest feature is what the agent refuses to claim

Public evidence is not internal truth

Management statements remain attributable claims. Missing workflow data stays missing.

No invented ROI

The agent does not fabricate costs, returns, implementation effort, data readiness, or causal impact.

No autonomous high-impact action

A person approves every consequential decision, and pilots begin in shadow mode.

No broad transformation leap

The system recommends one bounded test, not a companywide roadmap from public documents.

No hidden uncertainty

Facts, inference, assumptions, counter-signals, and source gaps remain separate.

Missing evidence is treated as an output rather than an inconvenience to hide. This claim discipline is part of the product architecture, not a disclaimer added after the recommendation.

How the specialist agent is structured

Hermes profile	Persistent specialist role, model configuration, tools, memory boundary, and Bot Mode identity.
Custom system persona	Standing evidence standards, research behavior, scoring rules, and deliverable requirements.
Domain skill	Reusable source hierarchy, evidence-ledger schema, problem-selection rules, scoring method, and pilot contract.
Research tools	Web, browser, and document extraction for investor materials, SEC filings, presentations, and transcripts.
Code execution	Deterministic scoring, arithmetic, consistency checks, and sensitivity analysis.
Document layer	Structured Markdown, editable DOCX, and shareable PDF executive memos.

Sources become evidence records. Evidence supports problems. Problems produce bounded opportunities. Scoring selects one pilot. A human approves or stops it.

What is complete and what still needs proof

Completed • Working specialist profile and domain skill. • Evidence hierarchy and quotation policy. • Transparent scoring and sensitivity warnings. • 30-day pilot contract and memo template. • Full PepsiCo and JPMorgan Chase demonstrations. • PDF, DOCX, Markdown, and evidence-ledger artifacts. • Public video demonstration and GitHub repository.	Still to validate • Practitioner review of evidence quality and usefulness. • A human-reviewed benchmark set across sectors. • Recommendation consistency across repeated runs. • Structured SEC ingestion and wider transcript coverage. • Measured adoption, ROI, time savings, or production impact.
--	---

Next, domain practitioners should score the case studies using a written rubric. The evaluation should track quotation accuracy, unsupported-claim rate, ranking stability, pilot completeness, and reviewer disagreement. A third-sector case should publish what changed after review, not merely add another polished memo.

Challenge the evidence, weights, and stop condition

The useful question is not whether the interface looks polished. It is whether the evidence supports the problem and whether the pilot could produce a credible go-or-stop decision.



Watch the video demonstration
youtu.be/HQ1QvQGP5NI



Explore the repository and case studies
github.com/vedantjainsc-hub/earnings-ai-strategist

Questions for reviewers

- Does each operating problem follow from the cited evidence?
- Which assumption would you validate first inside the company?
- Would different scoring weights change the recommended pilot?
- Are the baseline, target, human approval point, and stop condition decision-ready?
- What would make you reject the pilot before Day 6?

Disclaimer. This project is an independent demonstration using public information. It is not affiliated with or endorsed by PepsiCo, JPMorgan Chase, or any company analyzed by the agent. It is not investment, legal, accounting, or financial advice. The case-study recommendations are hypotheses for internal validation, not claims about current company operations or production readiness.