

JPMorgan Chase & Co. — Q2 2026 Earnings-to-AI Strategy Decision Memo

Ticker: NYSE: JPM Fiscal period: Second quarter ended June 30, 2026 As of: September 22, 2026 Name resolution: The user's "JP Morgan & Co" is treated as JPMorgan Chase & Co. Purpose: Internal operating decision support—not investment advice or equity research.

1. Decision requested

Approve a 30-day, shadow-mode pilot of an AI use-case evidence and portfolio review assistant for one non-customer-facing operations portfolio. The assistant would assemble traceable evidence for a monthly human decision: advance, hold, merge, or retire each selected AI use case. It would not approve funding, deploy models, change controls, make employment decisions, or communicate with clients.

2. Executive summary

- Observed operating signal: Jamie Dimon said the Firm had almost 1,000 AI use cases, while only about 50 were “really important,” spanning risk, fraud, marketing, hedging, prospecting, note-taking, idea generation, and document reading.[3]
- Recommended opportunity: An evidence-grounded portfolio review assistant that normalizes use-case records, links every benefit and control claim to current evidence, identifies apparent duplication, and exposes missing owners, metrics, or approvals for human disposition.
- Why now: The Firm is investing broadly in AI and technology while management emphasizes both customer benefit and workforce transition. A bounded governance workflow can test whether review quality and speed improve without delegating capital, risk, or people decisions to a model.[1][3][4]
- **Confidence:** Medium. Public evidence establishes portfolio scale and management priorities, but not the quality of the Firm's existing AI inventory, review cadence, internal controls, or data accessibility.
- Decision posture: Validate before building. If Days 1–5 show that existing governance already provides timely, complete, evidence-linked portfolio reviews, stop rather than duplicate it.

3. Company and reporting-period scope

JPMorgan Chase & Co. (NYSE: JPM) is the exact company analyzed. This memo covers the latest completed reporting period available as of the stated date: Q2 2026, ended June 30, 2026.

The Firm released its earnings release, presentation, and company-published transcript on July 14, 2026.[1][2][3]

Its Form 10-Q for the period was filed on August 6, 2026.[4][5]

The analysis is limited to operational evidence concerning AI portfolio governance, customer-outcome measurement, and workforce transition. It does not assess the Firm's stock, valuation, capital adequacy, creditworthiness, or investment merits.

4. Source manifest

ID	Source	Period	Published / filed	Type	Accessed	URL
S1 [1]	"2Q26 Earnings Press Release"	Quarter ended June 30, 2026	July 14, 2026	Company earnings release; primary	September 22, 2026	Direct PDF
S2 [2]	"2Q26 Earnings Presentation"	Quarter ended June 30, 2026	July 14, 2026	Company investor presentation; primary	September 22, 2026	Direct PDF
S3 [3]	"2Q26 Earnings Transcript"	Quarter ended June 30, 2026	July 14, 2026	Company-published transcript; primary	September 22, 2026	Direct PDF
S4 [4]	Form 10-Q for the quarterly period ended June 30, 2026	Quarter ended June 30, 2026	August 6, 2026	SEC filing; primary	September 22, 2026	Company-hosted filing PDF

ID	Source	Period	Published / filed	Type	Accessed	URL
S5 [5]	EDGAR filing index, accession 0001628280-26-054343	Quarter ended June 30, 2026	August 6, 2026	SEC filing index; primary	September 22, 2026	EDGAR

Source completeness: No required source gap was identified. JPMorganChase published the release, presentation, and transcript; the SEC provides the corresponding 10-Q index and filing. Management statements remain management claims rather than independently verified facts.

5. Three operational problems with exact quotations

Problem 1 — A very large AI portfolio requires disciplined evidence-based prioritization

Fact: Dimon said, “We have a lot of NPVs that we know we have. The whole company is working this at this point. And I think there’s almost 1,000 use cases today, though I would say that the really important ones are 50 – probably 50 across risk, fraud, marketing, hedging, prospecting, note-taking, idea generation, document reading.” — Jamie Dimon, Chairman and CEO, Q&A, company earnings transcript, July 14, 2026, PDF page 12.[3]

Corroborating fact: The 10-Q reported noninterest expense of \$27.3 billion, up 15%, including continued investment in marketing and technology; the filing also attributed the increase to compensation, brokerage and distribution, and occupancy factors.[4]

Inference: With roughly 1,000 use cases and a much smaller set described as especially important, portfolio reviewers face a recurring classification and evidence-normalization task: distinguish strategic use cases from experiments, duplicates, stale records, and low-evidence claims. This does not imply that current governance is weak.

Assumption requiring validation: The Firm maintains a reviewable registry with stable use-case IDs, owners, lifecycle states, costs, approved benefit definitions, control evidence, adoption data, and links to authoritative artifacts.

Counter-evidence / limitation: Management said it already knows many use-case NPVs and analyzes AI continuously.[3] Existing portfolio governance may already be mature; the pilot should proceed only if a measured review bottleneck remains.

Confidence: High that portfolio breadth exists; low-to-medium that evidence assembly is the material constraint.

Problem 2 — AI value must be demonstrated through customer outcomes, not assumed margin expansion

Fact: Dimon said, “The ultimate beneficiary of AI will be our customers. And in a competitive capitalist world, we all will use AI to do a better job for the customers, and we can’t just say, oh, it’s going to increase our margins and we’re going to keep that.” — Jamie Dimon, Chairman and CEO, Q&A, company earnings transcript, July 14, 2026, PDF page 7.[3]

Inference: Portfolio decisions need consistent customer- or user-outcome evidence—such as lower error rates, faster resolution, improved control quality, or better service—not only projected cost or margin benefits. The operating problem is comparable instrumentation and evidence, not a claim that customer outcomes are absent.

Assumption requiring validation: Each in-scope use case can be linked to a defined beneficiary, baseline, numerator and denominator, guardrail metrics, and an approved observation window.

Counter-evidence / limitation: Dimon also said the Firm has NPVs it knows it has, and the earnings materials reported strong financial performance.[1][3] Public sources do not establish whether customer-outcome measurement is missing, inconsistent, or already embedded in governance.

Confidence: Medium on the need for outcome discipline; low on the size of any internal measurement gap.

Problem 3 — AI-driven efficiency creates retraining and redeployment work that must remain human-controlled

Fact: Dimon said, “We are preparing to make sure we can retrain our people, and we have had discrete areas where we did reduce jobs by 30% or 40%. And most of those people were offered jobs elsewhere.” — Jamie Dimon, Chairman and CEO, Q&A, company earnings transcript, July 14, 2026, PDF page 12.[3]

Inference: In affected units, managers and HR teams may need to map employees’ validated skills and preferences to approved roles, training, and transition pathways. This is an evidence and workflow coordination problem, not permission for automated employment decisions.

Assumption requiring validation: The Firm has lawful, consented, current skills and vacancy data, documented mobility rules, human appeal routes, and fairness monitoring appropriate to each jurisdiction.

Counter-evidence / limitation: Management said most affected people in the cited discrete areas were offered jobs elsewhere, indicating that redeployment processes already exist.[3] The public statement supplies no population size, geography, role mix, acceptance rate, or outcome denominator.

Confidence: Medium that transition work exists; low on whether AI assistance would improve it safely.

6. Ranked AI opportunities and transparent scoring

Formula: $0.35 \times \text{Impact} + 0.25 \times \text{Feasibility} + 0.25 \times \text{Data readiness} + 0.15 \times (6 - \text{Risk})$ Scale: 1 = low, 5 = high. Risk is 1 = low and 5 = high; safety equals 6 - Risk. All scores are provisional because no public source establishes internal system architecture, data quality, governance maturity, or current workflow performance.

Rank	Opportunity and bounded decision	Linked problem	Impact	Feasibility	Data readiness	Risk	Safety	Weighted score
1	AI use-case evidence and portfolio review assistant — a portfolio council decides advance / hold / merge / retire	1 and 2	5	4	3	2	4	4.10 / 5
2	Customer-outcome instrumentation exception monitor — a product owner decides which use cases require measurement remediation	2	4	3	3	3	3	3.35 / 5
3	Workforce redeployment evidence navigator — an authorized HR partner decides which mobility options to discuss with an employee	3	4	3	3	4	2	3.20 / 5

Scoring rationale and controls

1. AI use-case evidence and portfolio review assistant

- Impact 5: Targets the disclosed gap between nearly 1,000 use cases and roughly 50 described as especially important.
- Feasibility 4: Document retrieval, schema validation, exact evidence extraction, duplicate-candidate clustering, and missing-field detection can be bounded and run in shadow mode.
- Data readiness 3 (provisional): A portfolio registry likely exists, but data completeness, permissions, identifier consistency, and evidence links are unknown.
- Risk 2 / safety 4: The assistant prepares a packet only; authorized humans retain funding, lifecycle, control, and deployment decisions.
- Non-AI baseline: Current portfolio spreadsheet, registry, or governance-deck preparation process plus deterministic completeness rules.
- Principal failure mode: The assistant makes an incomplete or stale use case appear comparable and decision-ready.

2. Customer-outcome instrumentation exception monitor

- Impact 4: Supports the management principle that customers should benefit, and may improve scale/stop decisions.
- Feasibility 3: Outcome definitions vary materially across risk, fraud, marketing, operations, and employee workflows.
- Data readiness 3 (provisional): Telemetry may be distributed across products and control systems with different denominators.
- Risk 3 / safety 3: Poorly chosen proxies could reward activity while harming service, fairness, or control outcomes.
- Non-AI baseline: Metric dictionary, dashboard thresholds, and periodic product/control review.
- Principal failure mode: Optimizing an incomplete proxy or comparing non-equivalent populations.

3. Workforce redeployment evidence navigator

- Impact 4: Could reduce manual search and make approved mobility options easier to review.
- Feasibility 3: Skills extraction and vacancy matching are technically feasible but organizationally complex.
- Data readiness 3 (provisional): Skills, performance, preferences, eligibility, and vacancies may be incomplete or jurisdiction-specific.

- Risk 4 / safety 2: Employment, privacy, fairness, accessibility, labor, and adverse-impact risks are high.
- Non-AI baseline: HR-led internal mobility process and employee-directed job search.
- Principal failure mode: Biased or incomplete recommendations narrow a person's opportunities.

Sensitivity warning: The top option leads by 0.75 points, but the score is not robust to materially weaker evidence. If its impact falls from 5 to 4 and data readiness falls from 3 to 1, its score declines by 0.85 to 3.25, below opportunity 2. The ranking therefore remains provisional until the registry, governance process, and baseline are inspected.

7. Recommended 30-day pilot

Pilot contract

Name: AI Use-Case Evidence Review Assistant — shadow pilot Decision owner role: Firmwide or line-of-business executive accountable for AI portfolio governance (exact title to validate internally) Frontline users: AI portfolio manager, use-case sponsor, control partner, finance/benefits reviewer, and governance secretary Recurring decision: At one monthly review, decide whether each selected use case should advance, hold, merge, or retire, based on authoritative evidence and recorded human rationale.

In scope

- One non-customer-facing operations portfolio; document-reading or note-taking use cases are preferred if approved because they were explicitly named and can avoid direct credit, fraud, trading, or employment decisions.[3]
- An assumed 15–25 use cases, validated on Day 1.
- Read-only ingestion of approved use-case records, business cases, control status, model inventory links, adoption telemetry, incident records, approved benefit measurements, and decision history.
- Controlled taxonomy for owner, beneficiary, workflow, stage, model/data dependency, control tier, approved metrics, evidence freshness, and decision state.
- Exact evidence spans, missing-evidence flags, stale-record warnings, apparent duplicate clusters, and a human-editable review packet.

Out of scope

- No funding approval, model approval, production deployment, access change, control waiver, client communication, workforce action, or automatic lifecycle state change.
- No new ROI forecast generated by a language model.
- No use of customer transaction content, employee performance data, or production prompts/outputs unless separately approved and necessary.
- No cross-line-of-business ranking using incomparable benefit definitions.
- No replacement of model risk management, cybersecurity, privacy, legal, compliance, audit, or business ownership.

Required internal data and approvals

- Authoritative use-case registry and stable IDs.
- Current owner, sponsor, lifecycle state, target user/beneficiary, workflow, investment-to-date where approved, and recurring cost where approved.
- Benefit definition, baseline provenance, observation window, numerator/denominator, actual result, and finance/benefit-owner sign-off.
- Model, data, security, privacy, legal/compliance, operational risk, and control approval status—not underlying restricted content unless required.
- Adoption, reliability, incident, override, complaint, and guardrail telemetry appropriate to the use case.
- Security, data-classification, retention, model/vendor, and least-privilege approvals for the pilot environment.

Baseline and day-30 targets

Baseline, measured in Days 1–5

1. Median preparation minutes per use-case review packet: total hands-on preparation minutes ÷ completed packets, with the median reported across cases.
2. Evidence completeness: required fields with current authoritative evidence ÷ all required fields.
3. Traceability: factual claims linked to an accessible artifact and version ÷ factual claims reviewed.
4. Review rework: packets returned for missing or inconsistent evidence ÷ packets submitted.
5. Duplicate-screen coverage: cases assessed under the approved similarity rubric ÷ in-scope cases.

Proposed approval targets—not forecasts

- At least 30% lower median packet-preparation time versus the same team's manual baseline.
- At least 20 percentage points higher evidence completeness, or ≥95% if the baseline is already high.
- 100% traceability for factual claims and benefit measurements included in the packet.
- 100% duplicate-screen coverage, with every candidate pair reviewed by humans; the target is coverage and decision quality, not a predetermined number of mergers.
- At least 70% reviewer agreement that assistant-raised missing-evidence and duplicate flags are actionable or correctly abstained.
- Zero autonomous portfolio, funding, control, production, client, or employment actions.

Weekly plan

Period	Work	Evidence produced	Human approval point
Days 1–5 — Validate	Map the current review; compare existing tools; define one portfolio and decision; inventory data; freeze taxonomy and required evidence; establish baseline; complete security/privacy gate	Current-state workflow, problem/data card, access matrix, baseline, non-AI rules baseline	Owner and control partners approve or stop before ingestion
Days 6–12 — Build	Use a frozen export or read-only connectors; implement schema and freshness checks; retrieve exact evidence; create duplicate-candidate logic; add abstention and audit fields	Reproducible prototype, lineage record, frozen rubric, 10-case test set	Portfolio manager and control partner approve rubric and representative cases
Days 13–21 — Test	Blind-compare against historical/current packets; classify omissions, stale links, false duplicate flags, inconsistent metrics, and permission failures; red-team prompt injection and cross-domain leakage	Error taxonomy, actionability score, access-control results, counterexamples	Two reviewers adjudicate disagreements; owner approves shadow review only if guardrails pass
Days 22–30 — Shadow pilot	Run alongside the existing review; humans use the current approval path; measure preparation time, completeness, rework, traceability, and adoption	Paired scorecard, audit log, user feedback, stop/extend/proceed recommendation	Day-30 steering review chooses stop, extend, or controlled production phase

Human approval and failure behavior

- The assistant displays source, version, date, exact evidence span, and confidence/abstention state for every factual assertion.
- Missing, inaccessible, conflicting, or stale evidence remains explicitly unresolved; the assistant must not infer completion.
- Similarity output is labeled candidate duplication, never proof that two use cases should merge.
- Benefit comparisons are blocked when baselines, populations, time windows, or denominators differ materially.
- Only authorized reviewers can choose advance, hold, merge, or retire; rationale and evidence snapshot are recorded.
- Any control, privacy, legal/compliance, operational risk, cybersecurity, or model-risk objection routes to the responsible human authority.

Stop conditions

Stop immediately for unauthorized access, cross-domain or client-data leakage, an automatic state/funding/control change, or a fabricated evidence citation. Stop or redesign if: (1) more than 20% of audited packets contain a material mismatch with the authoritative registry in either of two consecutive samples; (2) traceability falls below 100%; (3) the baseline cannot be reproduced; (4) existing governance already meets the targets at lower complexity; or (5) the Day-30 comparison shows no meaningful improvement in preparation time, evidence completeness, or rework.

Measurement method

Use paired cases: the same in-scope use cases receive the current manual packet and the shadow assistant packet. Capture hands-on preparation time from task start to review-ready packet. Score completeness against the frozen rubric, not a model-generated rubric. Two authorized reviewers independently label missing-evidence and duplicate-candidate flags; report agreement, disagreements, abstentions, sample size, median and range, rework, and all guardrail incidents. Do not extrapolate a 30-day workflow result into Firmwide savings or ROI.

8. Risks, assumptions, missing internal evidence, and non-goals

Risks and guardrails

- Sensitive portfolio information: least-privilege access, approved retrieval scope, no training on internal content, and immutable access logs.
- Stale or persuasive misinformation: exact-source spans, version/freshness checks, authoritative-source allowlist, and mandatory abstention.
- Metric comparability: enforce definitions, denominator and baseline provenance, and block invalid cross-use-case comparisons.
- Automation bias: display counter-evidence, unresolved fields, and the non-AI baseline; require human rationale for every decision.
- Control bypass: consume approval status read-only; never infer or issue a control approval.
- Duplicate false positives: use similarity only to open a review; require owner and architecture/control confirmation before any merge decision.
- Portfolio bias: do not privilege use cases merely because evidence is easier to collect; show missingness separately from strategic value.

Assumptions to validate

- A single registry or reconciled view exists for the selected portfolio.
- Stable use-case IDs connect decisions, control records, metrics, and evidence.
- The selected portfolio has enough review activity for a 30-day comparison.
- Existing governance permits a frozen, read-only pilot dataset.
- Reviewers agree on required fields, evidence freshness, and decision definitions.
- A non-customer-facing portfolio can be selected without weakening the relevance of the test.

Missing internal evidence

- Current number of active versus experimental, merged, retired, or stale use cases.
- Existing AI inventory, governance tooling, review frequency, and packet-preparation effort.
- Percentage of use cases with approved baselines, measured outcomes, and current control evidence.
- Duplication, rework, late control findings, deployment failure, incident, and abandonment rates.
- Internal cost, model/vendor architecture, permissions, and retention constraints.
- Whether “almost 1,000” refers to ideas, prototypes, deployed systems, or another lifecycle definition.

Non-goals

This memo does not recommend reducing the AI portfolio to 50, ranking employees, automating employment decisions, estimating Firmwide AI ROI, bypassing model risk management, or replacing existing governance. It does not infer that 1,000 use cases are excessive, that duplication exists, or that current portfolio controls are deficient.

9. Executive approval checklist

- [] Confirm the exact decision owner and one non-customer-facing portfolio.

- [] Confirm the pilot does not duplicate existing registry or governance capabilities.
- [] Approve the use-case taxonomy, evidence rubric, decision states, and authoritative sources.
- [] Approve security, privacy, legal/compliance, operational risk, model/vendor, and retention controls.
- [] Confirm the baseline sample and metric definitions before building.
- [] Keep funding, control, production, client, and people decisions with authorized humans.
- [] Treat proposed Day-30 thresholds as evaluation criteria—not forecasts.
- [] Authorize immediate stop for leakage, fabricated evidence, or autonomous state changes.
- [] Name the Day-30 steering group and decision rights: stop, extend, or controlled next phase.

10. Source notes with URLs

Sources 1–5 are company or SEC primary materials. The URL mapping below is generated from the citation ledger.

Sources

[1] <https://jpmorganchase.com/content/dam/jpmc/jpmorgan-chase-and-co/investor-relations/documents/quarterly-earnings/2026/2nd-quarter/6cded9fd-a164-4e6c-8cff-377357cf105c.pdf> — JPMorganChase 2Q26 Earnings Press Release

[2] <https://jpmorganchase.com/content/dam/jpmc/jpmorgan-chase-and-co/investor-relations/documents/quarterly-earnings/2026/2nd-quarter/9e8f3cef-bd2a-4140-b50f-83c4fba3f23c.pdf> — JPMorganChase 2Q26 Earnings Presentation

[3] <https://jpmorganchase.com/content/dam/jpmc/jpmorgan-chase-and-co/investor-relations/documents/quarterly-earnings/2026/2nd-quarter/2Q26-earnings-transcript.pdf> — JPMorganChase 2Q26 Earnings Transcript

[4] <https://jpmorganchase.com/content/dam/jpmc/jpmorgan-chase-and-co/investor-relations/documents/quarterly-earnings/2026/2nd-quarter/corp-q2-2026.pdf> — JPMorgan Chase & Co. Q2 2026 Form 10-Q

[5] <https://www.sec.gov/Archives/edgar/data/19617/000162828026054343/0001628280-26-054343-index.html> — SEC Form 10-Q Filing Index

Evidence appendix

Claim ID	P1-E1
Problem	AI portfolio prioritization
Exact quotation	"We have a lot of NPVs that we know we have. The whole company is working this at this point. And I think there's almost 1,000 use cases today, though I would say that the really important ones are 50 – probably 50 across risk, fraud, marketing, hedging, prospecting, note-taking, idea generation, document reading."
Speaker / anchor	Jamie Dimon, Q&A, transcript p. 12 [3]
Interpretation	Portfolio breadth creates a recurring review burden
Limitation	Does not define lifecycle stage or imply excess use cases
Claim ID	P1-E2
Problem	Investment context
Exact quotation	"Noninterest expense was \$27.3 billion, up 15%, predominantly driven by higher compensation expense as a result of higher revenue-related compensation, wage inflation and growth in the number of employees, as well as higher brokerage expense and distribution fees, continued investments in marketing and technology, and higher occupancy expense."
Speaker / anchor	Form 10-Q, Firmwide overview, p. 6 [4]
Interpretation	Selectivity matters in a broad investment environment
Limitation	Expense growth had multiple revenue-linked causes
Claim ID	P2-E1
Problem	Customer outcome principle
Exact quotation	"The ultimate beneficiary of AI will be our customers. And in a competitive capitalist world, we all will use AI to do a better job for the customers, and we can't just say, oh, it's going to increase our margins and we're going to keep that."
Speaker / anchor	Jamie Dimon, Q&A, transcript p. 7 [3]
Interpretation	Portfolio evidence should include customer/user outcomes
Limitation	Does not establish that current measurement is deficient
Claim ID	P3-E1
Problem	Retraining and redeployment
Exact quotation	"We are preparing to make sure we can retrain our people, and we have had discrete areas where we did reduce jobs by 30% or 40%. And most of those people were offered jobs elsewhere."
Speaker / anchor	Jamie Dimon, Q&A, transcript p. 12 [3]
Interpretation	Transition workflows need coordinated evidence and human oversight
Limitation	No denominator, geography, role mix, or outcome rate